Stochastic/probabilistic/Monte Carlo methods

A large number of physical and other processes in nature are random, particularly those quantum mechanical in nature like radioactive decay and other quantum mechanical processes

Some processes are not actually completely random but for our purposes are essentially random: flipping a coin, rolling dice, spinning a roulette wheel, Brownian motion

Monte Carlo simulation

https://en.wikipedia.org/wiki/Monte_Carlo_method#History

In the late 1940s, Stanislaw Ulam invented the modern version of the Markov Chain Monte Carlo method while he was working on nuclear weapons projects at the Los Alamos National Laboratory. Immediately after Ulam's breakthrough, John von Neumann understood its importance. Von Neumann programmed the ENIAC computer to perform Monte Carlo calculations. In 1946, nuclear weapons physicists at Los Alamos were investigating neutron diffusion in fissionable material.[12] Despite having most of the necessary data, such as the average distance a neutron would travel in a substance before it collided with an atomic nucleus and how much energy the neutron was likely to give off following a collision, the Los Alamos physicists were unable to solve the problem using conventional, deterministic mathematical methods. Ulam proposed using random experiments. He recounts his inspiration as follows:

> The first thoughts and attempts I made to practice [the Monte Carlo Method] were suggested by a question which occurred to me in 1946 as I was convalescing from an illness and playing solitaires. The question was what are the chances that a Canfield solitaire laid out with 52 cards will come out successfully? After spending a lot of time trying to estimate them by pure combinatorial calculations, I wondered whether a more practical method than "abstract thinking" might not be to lay it out say one hundred times and simply observe and count the number of successful plays. This was already possible to envisage with the beginning of the new era of fast computers, and I immediately thought of problems of neutron diffusion and other questions of mathematical physics, and more generally how to change processes described by certain differential equations into an equivalent form interpretable as a succession of random operations. Later [in 1946], I described the idea to John von Neumann, and we began to plan actual calculations.[13]

Being secret, the work of von Neumann and Ulam required a code name.[14] A colleague of von Neumann and Ulam, Nicholas Metropolis, suggested using the name *Monte Carlo*, which refers to the Monte Carlo Casino in Monaco where Ulam's uncle would borrow money from relatives to gamble.[12] Using lists of "truly random" random numbers was extremely slow, but von Neumann developed a way to calculate pseudorandom numbers, using the middle-square method. Though this method has been criticized as crude, von Neumann was aware of this: he justified it as being faster than any other method at his disposal, and also noted that when it went awry it did so obviously, unlike methods that could be subtly incorrect.[15]

Actually, where does the name come from?

My research is very much a random process

Remember what I call that?

The dinosaur plot

Decays, too

Higgs boson couples proportional to mass of particle. But its decays are random and probabilistic!

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

Why Monte Carlo?

Monte Carlo assumes the system is described by probability density functions (PDF) which can be modeled. It does not need to write down and solve equation analytically/numerically.

PDF comes from
- Data driven
- Theory driven
- Data + Theory fitting

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

Particle physics uses MC for

(1) Detector design and optimization
- Complicate and huge detector
- Very expensive

(2) Simulation of particle interactions with detector's material

(3) Physics analysis
- New predicted physics: SUSY, UED, ...
- Event selection
- Background estimation
- Efficiencies of detector/algorithm/...

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

Monte Carlo Simulation in HEP

Choose model, constraints, parameters, decay chain of interest

Proposed Theory

Kinematics, information from a known (detectable) particles

Generator

Detector Simulation
- Hardware
- Software

Simulation Digitization Triggering

Offline software
- Event selection

Reconstruction Analysis

Experiment Triggering

Results Improvement

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

Monte Carlo generators

Generator (study @ parton level)

↓ ↓

Theoretical Model | **Shower generator**

Input: Model parameters.
Output: Four-vector of momenta of stable/quasi-stable particles produced in interactions

Example MC generator Inteface with CMSSW

Pythia6 ★	Phantom
Herwig6 ★	Hydjet
Pythia8 ★	Pyquen
ThePEG (Herwig++, Ariadne 5) ★	Cosmic Muon Generator
ALPGEN	Beam Halo Muon Generator
MadGraph	ExHuME
MC@NLO	Pomwig
POWHEG	BcGenerator
SHERPA	HARDCOL

A list of MC generator can be found in [http://www.hepforge.org/]

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

Monte Carlo event generator process

(1) Hard process: What do you want to study?
(2) Parton-shower phase:
(3) Hard particles decay before hadronizing: e.g. top, SUSY
(4) Hadronization: form observed hadron
(5) Unstable hadrons decay: Experimentally measured BR, phase-space distribution of the decay product

Can have loops and many higher-order corrections of a variety of types, too!

Spin Correlations in Monte Carlo Simulations

Peter Richardson

Cavendish Laboratory, University of Cambridge, Madingley Road, Cambridge, CB3 0HE, UK, and
DAMTP, University of Cambridge, Centre for Mathematical Sciences, Wilberforce Road, Cambridge, CB3 0WA, UK.

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

How does MC work in detector simulation?

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

How does MC work in detector simulation?

In electron case:

- ☑ With probability **p1**: ionize the gas, loose some momentum, produce N secondary electrons with momentum p_{e1N1},...
- ☑ Do nothing with probability **1-p1**
- ☑ Generate random number **r** in the range [0,1]
- ☑ If **r<p1**, generate momenta of secondary electrons, add them to particles list, reduce the momentum of initial electron.

More on particle physics (nice slides by Srimanobhas)

https://indico.cern.ch/event/92209/contributions/2114409/attachments/1098701/1567290/CST2010-MC.pdf

How does MC work in detector simulation?

In photon case:

- ☑ With probability **p1**: convert and produce electron-positron pair
- ☑ With probability **p2**: Compton scattering
- ☑ With the probability **p3**: ionizing the matter
- ☑ Generate random number **r** in the range [0,1]
- ☑ Three cases:
  - ☑ **r < p1**
  - ☑ **p1 < r <p1+p2**
  - ☑ **p1+p2 < r < p1+p2+p3**

Not as simple as it seems

Take 100 electrons (not so many!) and allow them to exist in spin-up or spin-down states. What is the total number of available configurations for the system?

$2^{100} = 1.3 \times 10^{30}$ (ooof)

If we want to study this system, we will need to be clever about what random configurations we examine

One inherent problem

By their very nature, computers are NOT random. They are deterministic. That is a very good thing - I don't want to work on a computer that randomly produces output. But if we then want to generate random numbers, we have a problem! We can connect our computer to a Geiger counter, but that isn't super ideal. Instead we use pseudo-random numbers, ie numbers that appear random as far as we can tell, but in fact are not random

"Any one who consider arithmetical methods of producing random digits is, of course, in a state of sin." John von Neumann, 1951

What do we want in our random numbers?

Easiest place to start, probability function

$p(x) = 1$ if $0 \leq x \leq 1$, 0 otherwise

- **Uniformity**
  - Don't want gaps where numbers are not picked, this isn't random
- **Independence**
  - Don't want this random number to appear to depend on the previous one.
- **Performance of algorithm**
  - Needs to be efficient, fast and something that can be parallelized, ideally
- **Replicability**
  - Need to be able to reproduce results. This is important for debugging
- **Long length cycle**
  - At some point, all algorithms that purely run on a computer will cycle and start all over. Don't want that to happen for a long time

Needed for numerical algorithms, simulations, Monte Carlo methods, encryption……

My favorite application of random numbers

```
[20] from matplotlib import pyplot as plot
     from numpy import random,sqrt

     N = 1000000
     xs_in=[]
     ys_in=[]
     xs_out=[]
     ys_out=[]
     for i in range(N):
       x = random.uniform(-1,1,1)
       y = random.uniform(-1,1,1)
       r = sqrt(x*x+y*y)
       if (r < 1):
         xs_in.append(x)
         ys_in.append(y)
       else:
         xs_out.append(x)
         ys_out.append(y)

     plot.axis([-1,1,-1,1])
     plot.gca().set_aspect('equal', adjustable='box') ### make the plot square
     plot.scatter(xs_in,ys_in,color='r',s=0.0001) ## make size of markers small for so many points
     plot.scatter(xs_out,ys_out,color='b',s=0.0001)  ## make size of markers small for so many points
     plot.show()
     fraction_in = len(xs_in)/N
     ## Area of Rectangle is 2*2 = 4
     ## Area in circle is pi*r*r = pi*1*1 = pi
     ## So fraction inside is fraction in circle = pi/4 and pi = 4*fraction
     print("Fraction in a circle of radius 1 = ", fraction_in, "and so pi = ", 4*fraction_in)
```

```
Fraction in a circle of radius 1 =  0.784895 and so pi =  3.13958
```

Let’s look at this code together

Let's look at how to speed up this code

```
# Faster way to check if the points are inside the circle
# General rule is to avoid "for" loops if a simple numpy expression can be written

from numpy import zeros,random,pi
import time

tstart = time.perf_counter()

N = 5000000

results = zeros(N)
x_values = random.uniform(-1, 1, N)
y_values = random.uniform(-1, 1, N)
# the following means: "for the indices where the point satisifies x**2 + y**2 <= 1, set results to 1"
# note that we don't actually need the square root, either! And the mean is perfect to use since for zeros or ones it is just the fraction passing, which is what we want
results[x_values**2 + y_values**2 <= 1] = 1

tend = time.perf_counter()

print("Estimate for pi with N = ",N," is ",4*results.mean())
print("The estimate took ",(tend - tstart),"seconds")
```

```
Estimate for pi with N =  5000000  is  3.1418008
The estimate took  0.1766990120000287 seconds
```

Let's look at this code together and then discuss it

Slightly even faster

```python
# Even faster - skip filling the results array, since we don't actually care *which* points are inside the
# circle, just how many there are
# when the expression is True, sum counts it as 1, and when the expression is False, sum counts it as zero

from numpy import zeros,random,pi,sum
import time

tstart = time.perf_counter()

N = 5000000

results = zeros(N)
x_values = random.uniform(-1, 1, N)
y_values = random.uniform(-1, 1, N)
count = sum(x_values**2 + y_values**2 <= 1)

tend = time.perf_counter()
print("Estimate for pi with N = ",N," is ",4*count/N)
print("The estimate took ",(tend - tstart),"seconds")
```

```
Estimate for pi with N =  5000000  is  3.1427208
The estimate took  0.15959697000005235 seconds
```

Also a lot shorter than the original!

Slightly even faster


```
# Even faster - skip filling the results array, since we don't actually care *which* points are inside the
# circle, just how many there are
# when the expression is True, sum counts it as 1, and when the expression is False, sum counts it as zero
### now do a large scale check
from numpy import zeros,random,pi,sum
import time

tstart = time.perf_counter()

N = 50000000

results = zeros(N)
x_values = random.uniform(-1, 1, N)
y_values = random.uniform(-1, 1, N)
count = sum(x_values**2 + y_values**2 <= 1)

tend = time.perf_counter()
print("Estimate for pi with N = ",N," is ",4*count/N)
print("The estimate took ",(tend - tstart),"seconds")
```

```
Estimate for pi with N =  50000000  is  3.14116336
The estimate took  1.9479021709999955 seconds
```

Not bad!

On random strings (for the geeks in the audience)

Your PHYS410/510 instructor

How do you generate a random string? ...
Put a web designer in front of VIM
and tell him to save and exit.

My favorite application of random numbers

But what is the uncertainty on our method? The probability of being inside the circle = π/4 = p for a binomial.

Recall way back from statistics class that the variance on the sum of drawing from a binomial n times with probability of success p is np(1-p), so that the "1 sigma" uncertainty on it is sqrt(np(1-p)).

That is the total number the uncertainty on the fraction in the circle is then that number divided by n, or sqrt(p(1-p)/n).

Our estimate of pi is 4 times that value, so it has error 4*sqrt(p(1-p)/n).

In other words

```
from matplotlib import pyplot as plot
from numpy import random,sqrt,pi

### probability of being inside = pi/4 = p for a binomial
### variance of a binomial = n*p*(1-p), sigma is sqrt(n*p*(1-p)).
### But that sigma is on the total number, the sigma on the fraction is that
### number divided by n = sqrt(p*(1-p)/n)
### Value of pi = 4p, uncertainty on it is 4*sqrt(p*(1-p)/n)

N = 500000
xs = []
ys = []
ydiffs = []
binomialdiffs = []
nin = 0
p = pi/4
for i in range(N):
  x = random.uniform(-1,1,1)
  y = random.uniform(-1,1,1)
  r = sqrt(x*x+y*y)
  if (r < 1):
    nin = nin+1
  thispi = 4.*nin/(i+1) ### divide by i+1 because we finished ith entry already
  xs.append(i)
  ys.append(thispi)
  ydiffs.append(abs(pi-thispi))
  binomialdiffs.append(4*sqrt(p*(1-p)/(i+1)))

figs,axes = plot.subplots(2,figsize=(15,15))
axes[0].scatter(xs,ys,color='b')
axes[0].set(ylabel="Estimated value of pi")
axes[0].set(ylim=[0.98*pi,1.02*pi])
axes[1].scatter(xs,ydiffs,color='r')
axes[1].scatter(xs,binomialdiffs,color='g')
axes[1].set_title("trials")
axes[1].set(ylabel="Error on pi")
axes[1].set(ylim=[0,0.02])
plot.show()
```

Our estimate of pi is 4 times that value, so it has error 4*sqrt(p(1-p)/n), so our “error” scales as 1/sqrt(n). If we run 4x more trials, we only improve our accuracy by 50%

Tracking our results

Linear Congruential generator

Possibly the simplest pseudo-random number generator that you can think of implementing on a computer:

$$\mathbf{X_{n+1} = (aX_n + c) \bmod m}$$

Generates (n+1)th random number from the nth random number. The values a, c and m are constants, and X_0 is a seed (the first random number) or start value

Need to be careful about choices of a, c and m. And the random numbers will repeat after at most m values

Linear Congruential generator

```python
# Checking pseudo-random numbers - looks better
import matplotlib.pyplot as plt

N = 1000
a = 1234567
c = 123456789
m = 987654321
x = 1
y = 0
xvalues = []
yvalues = []

for i in range(N):
  y = (a*x+c)%m
  xvalues.append(x)
  yvalues.append(y)
  x = (a*y+c)%m
  xvalues.append(x)
  yvalues.append(y)

plt.scatter(xvalues,yvalues)
plt.show()
```

Not crazy!

Need to be careful about (psuedo!) random numbers

```
# Checking pseudo-random numbers
import matplotlib.pyplot as plt

N = 8
a = 5
c = 3
m = 8
x = 1
y = 0
xvalues = []
yvalues = []

for i in range(N):
  y = (a*x+c)%m
  xvalues.append(x)
  yvalues.append(y)
  x = (a*y+c)%m
  xvalues.append(x)
  yvalues.append(y)

plt.scatter(xvalues,yvalues)
plt.show()
```

Not so random! The points lie in a hyper-plane. Easy to see in small examples, but sometimes even fancy generators can have this (google RANDU generator for an example)

Brownian motion (Ex 10.3)

```
from math import log
from vpython import sphere,box,color,rate,vector
from random import randrange

L = 1001
N = 1000000
framerate = 1000

box(pos=vector(-L/2,0,0),length=1,height=L,width=1,color=color.green)
box(pos=vector(L/2,0,0),length=1,height=L,width=1,color=color.green)
box(pos=vector(0,-L/2,0),length=L,height=1,width=1,color=color.green)
box(pos=vector(0,L/2,0),length=L,height=1,width=1,color=color.green)
s = sphere(pos=vector(0,0,0), radius=5, color = color.white)

# Main loop
i = j = 0
for k in range(N):
  direction = randrange(4)
  if (direction == 0):
    if i < L/2: i += 1
  elif (direction == 1):
    if i > -L/2: i -= 1
  elif (direction == 2):
    if j < L/2: j += 1
  else:
    if j > -L/2: j -= 1
  rate(framerate)
  s.pos = vector(i,j,0)
```

Random motion of a particle (like a dust particle) in a gas. At each time step it bounces into another molecule in the gas (not shown) and chooses the direction to take. Let's watch together

Brownian motion (Ex 10.3), with matplotlib

```
#Brownian motion!

from math import log
from random import randrange
import matplotlib.pyplot as plt
from matplotlib import animation, rc
import numpy as np
from IPython.display import HTML

L = 201
N = 10000
####N = 1000000
framerate = 1

x = y = 0

def animate(i):
  global x,y
  # Main loop
  ###print(pos[i,0],pos[i,1])
  # Main loop
  direction = randrange(4)
  if (direction == 0):
    if x < L/2: x += 1
  elif (direction == 1):
    if x > -L/2: x -= 1
  elif (direction == 2):
    if y < L/2: y += 1
  else:
    if y > -L/2: y -= 1
  point.set_data(x,y)
  return point,

fig = plt.figure()
ax = plt.axes(xlim=(-L/2, L/2), ylim=(-L/2,L/2))
ax.set_aspect("equal")
# create a point in the axes
point, = ax.plot(x,y, marker="o")


ani = animation.FuncAnimation(fig, animate, interval=framerate, frames = N, blit=True)
###plt.show()
#ani.save('brownian_motion_2d.mp4', fps=60, extra_args=['-vcodec', 'libx264'])
rc('animation', html='jshtml')
ani
```

We can watch the output again (it's kind of mesmerizing, I think)

How to generate random numbers?

We examined the “uniform distribution” before. Normally this is a distribution between 0-1. What if we want a uniform number between a-b? We can simply rescale things:

$p(x) = 1$ if $0 \leq x \leq 1$, 0 otherwise
$q(x) = N*p(x)+z$ where $N = (b-a)$ and $z = a$

But what if we don’t want to generate numbers according to a Uniform distribution? What if we want numbers distributed according to a Poisson? An exponential? Otherwise?

How to generate random numbers?

Use the chain rule:

$P(y) = P(x)\ |dx/dy|$

Uniform distribution between a and b:

$P(x) = 1$ if $0 \leq x \leq 1$, 0 otherwise

$y(x) = a + (b-a)x$

So $x = (y-a)/(b-a)$ and $dx/dy = 1/(b-a)$

Then $P(y) = 1/|b-a|$ for $a < y < b$, 0 otherwise

How to generate random numbers?

Use the chain rule:

P(y) = P(x) ldx/dyl, remember that x is uniformly distributed between 0 and 1 and we want to relate x and y

$$P(y) = \mu e^{-\mu y}, 0 < y < \infty$$

$$P(y)dy = \mu e^{-\mu y} dy = P(x)dx$$

$$\int_0^y \mu e^{-\mu y'} dy' = \int_0^x P(x')dx'$$

$$\mu \int_0^y e^{-\mu y'} dy' = x$$

How to generate random numbers?

$$\mu \int_0^y e^{-\mu y'} dy' = x$$

$$1 - e^{-\mu y} = x$$

$$e^{-\mu y} = 1 - x$$

$$-\mu y = \ln(1 - x)$$

$$y = -\frac{1}{\mu} \ln(1 - x)$$

So if we generate x uniformly between 0 and 1 and plug into the above, then y will have the form we want

What if we want to generate Gaussian random numbers?

$$P(y) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{y^2}{2\sigma^2}}$$

$$\frac{1}{\sqrt{2\pi\sigma^2}} \int_{-\infty}^{y} e^{-\frac{y^2}{2\sigma^2}} dy = x$$

Unfortunately can't solve this, but let's try drawing two random numbers from Gaussians (with the same width)

$$P(y,z)dydz = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{y^2}{2\sigma^2}} \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{z^2}{2\sigma^2}} dydz$$

What if we want to generate Gaussian random numbers?

$$P(y,z)dydz = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{y^2}{2\sigma^2}} \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{z^2}{2\sigma^2}} dydz$$

$$P(y,z)dydz = \frac{1}{2\pi\sigma^2} e^{-\frac{y^2+z^2}{2\sigma^2}} dydz$$

Trick: Convert to Polar coordinates

$$P(y,z)dydz = rP(r,\theta)drd\theta$$

$$P(r,\theta)drd\theta = \frac{r}{\sigma^2} e^{-\frac{r^2}{2\sigma^2}} dr \frac{d\theta}{2\pi}$$

What if we want to generate Gaussian random numbers?

$$P(r,\theta)drd\theta = \frac{r}{\sigma^2}e^{-\frac{r^2}{2\sigma^2}}dr\frac{d\theta}{2\pi}$$

The angular distribution is trivial (uniform between 0 and $2\boldsymbol{\pi}$), and we know how to do it

$$P(r)dr = \frac{r}{\sigma^2}e^{-\frac{r^2}{2\sigma^2}}dr$$

$$\int P(r)dr = \int \frac{r}{\sigma^2}e^{-\frac{r^2}{2\sigma^2}}dr$$

$$q = -\frac{r^2}{2\sigma^2}, dq = \frac{-2r}{2\sigma^2}dr = \frac{-r}{2\sigma^2}dr$$

$$\int P(r)dr = \int \frac{r}{\sigma^2}(e^q)\frac{-\sigma^2}{r}dq = -\int e^q dq$$

What if we want to generate Gaussian random numbers?

$$\int_0^r P(r)dr = -\int_{q=0}^{q=-\frac{r^2}{2\sigma^2}} e^q dq \qquad q = -\frac{r^2}{2\sigma^2}$$

$$x = (1 - e^{-\frac{r^2}{2\sigma^2}})$$

$$(1-x) = e^{-\frac{r^2}{2\sigma^2}}$$

$$\ln(1-x) = -\frac{r^2}{2\sigma^2}$$

$$-2\sigma^2 \ln(1-x) = r^2$$

$$r = \sqrt{-2\sigma^2 \ln(1-x)}$$

x is Uniformly distributed, setting r like this and using a random value for θ lets us pick y and z! If we only want one Cartesian coordinate, we can throw away the other one

Rejection samples / accept-reject samples

There is another way to draw random numbers from any arbitrary distribution. It is called rejection samples (or accept-reject sampling).

Let's assume we want to know P(x) between x_0 and x_1 (it must be zero outside this range)

We also must know a number M that is larger than P(x) for any x (it doesn't have to be the smallest such number)

We generate x' uniformly between x_0 and x_1 and compute

$$y = P(x') / M$$

We then generate y' from a Uniform distribution between 0 and 1. If y' < y, accept x', otherwise reject x'

What is rejection sampling doing?

For each x' we randomly decide whether to keep the number depending on its probability y

Let's assume we want to know P(x) between x_0 and x_1 (it must be zero outside this range)

We also must know a number M that is larger than P(x) for any x (it doesn't have to be the smallest such number)

We generate x' uniformly between x_0 and x_1 and compute

$$y = P(x') / M$$

We then generate y' from a Uniform distribution between 0 and 1. If y' < y, accept x', otherwise reject x'

Example of rejection sampling

```python
import matplotlib.pyplot as plt

# Implementation of accept/reject sampling for a continuous variable.
# Pass the Python function, the range of potential x values as a tuple (xmin, xmax), and the maximum value for f(x) to assume
def accept_reject(func, rng, maxval):
    from random import uniform
    while True:
        xtest = uniform(*rng)
        y = func(xtest)/maxval
        if y > 1:
            print(f"Problem: function ({y*maxval}) has exceeded maxval {maxval} for x {xtest}")
        ytest = uniform(0,1)
        if ytest < y:
            return xtest


# example. some arbitrary PDF:
def quadratic(x):
    return 0.75*(1.-x*x) if -1 < x < 1 else 0

# some points:
points = []
N = 50000
for x in range(N):
  points.append(accept_reject(quadratic, (-1,1), 1))

# compare histogram of generated points to PDF
plt.hist(points, bins=20, density=True)
plt.plot(numpy.linspace(-1,1,1000), numpy.frompyfunc(quadratic, 1, 1)(numpy.linspace(-1,1,1000)))
plt.show()
```

Let’s check this

Another example

```python
import matplotlib.pyplot as plt
from numpy import exp,power,frompyfunc,linspace,sqrt,pi

# Implementation of accept/reject sampling for a continuous variable.
# Pass the Python function, the range of potential x values as a tuple (xmin, xmax), and the maximum value for f(x) to assume
def accept_reject(func, rng, maxval):
    from random import uniform
    while True:
        xtest = uniform(*rng)
        y = func(xtest)/maxval
        if y > 1:
            print(f"Problem: function ({y*maxval}) has exceeded maxval {maxval} for x {xtest}")
        ytest = uniform(0,1)
        if ytest < y:
            return xtest


xmax=5

# truncated Gaussian, width 1, mean 0
def trunc_gaus(x):
  return 1./(sqrt(2*pi))*exp(-0.5*power(x,2)) if abs(x) < xmax else 0

# some points:
points = []
N = 50000
for x in range(N):
  points.append(accept_reject(trunc_gaus, (-xmax,xmax), 0.5))

# compare histogram of generated points to PDF
plt.hist(points, bins=50,density=True)
plt.plot(linspace(-xmax,xmax,1000), frompyfunc(trunc_gaus, 1, 1)(linspace(-xmax,xmax,1000)))
plt.show()
```

Monte Carlo integration

```python
# Ugly function that we want to integrate
from math import sin
from random import random
from matplotlib import pyplot as plot
import numpy as np

def f(x):
  return (sin(x/(x*(2-x)*(3-x))))**4

epsilon = 1e-8 ### so we don't divide by zero
xmin=0+epsilon
xmax=2-epsilon
N=1000
xlist = np.linspace(xmin, xmax, N)
ylist=[]
for x in xlist:
  ylist.append(f(x))

plot.plot(xlist,ylist)
plot.show()
```

What is this integral? It's rapidly varying near the edges, so is there a good solution to estimating it?

$$I = \int_0^2 \sin^4 \left[\frac{1}{x(2-x)(3-x)} \right] dx$$

Monte Carlo integration

```
# Example MC integration
from math import sin
from random import random
from matplotlib import pyplot as plot
import numpy as np

def f(x):
  return (sin(x/(x*(2-x)*(3-x))))**4

epsilon = 1e-8 ### so we don't divide by zero
xmin=0+epsilon
xmax=2-epsilon
N=100000
xlist = np.linspace(xmin, xmax, N)
ylist=[]
count = 0
for i in range(N):
  randx = 2*random()
  randy = random()
  if (randy < f(randx)):
    count = count + 1

print(2*count/N)
```

```
0.41076
```

The integral is the “area under the curve”. And we know the total integral of the box from x=0 to x=2 (and y=0 to y=1) is 2*1 = 2

$$I = \int_0^2 \sin^4 \left[\frac{1}{x(2-x)(3-x)} \right] dx$$

How fast does this converge?

```
#Checking convergence of MC Integration
from math import sin
from random import random
from matplotlib import pyplot as plot
import numpy as np

def f(x):
  return (sin(x/(x*(2-x)*(3-x))))**4

epsilon = 1e-8 ### so we don't divide by zero
xmin=0+epsilon
xmax=2-epsilon
N=500000
xlist = []
ylist=[]
count = 0
for i in range(N):
  randx = 2*random()
  randy = random()
  if (randy < f(randx)):
    count = count + 1
  ylist.append(2*count/(i+1))
  xlist.append(i+1)

plot.figure(figsize=(10,10))
plot.ylim(0.35,0.45)
plot.scatter(xlist,ylist,color='b')
plot.ylabel("Estimated value of integral")
plot.xlabel("Trials")
plot.show()
```

Quite some variation, even after ~100,000 iterations! Slow convergence

How fast does this converge?

The odds of a trial passing are given by Binomial statistics, just like before when estimating pi, so the error scales exactly as 1/sqrt(N) again, which is quite slow convergence

```
[13] 1./(500000**0.5)
```

```
0.001414213562373095
```

Mean value method

$$I = \int_a^b f(x)dx$$

$$< f >= \frac{1}{b-a}\int_a^b f(x)dx = \frac{I}{b-a}$$

Rearranging: $$I = (b-a) < f >$$

But we have an easy method to estimate < f >. We just sample it at random points!

$$I = \frac{b-a}{N}\sum_{i=1}^{N} f(x_i)$$

Error on mean value method

$$< f >= \frac{1}{N} \sum_{i=1}^{N} f(x_i)$$

$$< f^2 >= \frac{1}{N} \sum_{i=1}^{N} [f(x_i)]^2$$

$$\mathrm{var} f =< f^2 > - < f >^2$$

$$I = (b-a) f$$

$$\mathrm{var} I = (b-a) \mathrm{var} f$$

Variance on I is proportional to 1/N, so uncertainty on I is proportional to 1/sqrt(N)

Comparison of methods

```python
# Mean Value vs MC Integration
from math import sin
from random import random
from matplotlib import pyplot as plot
import numpy as np

def f(x):
  return (sin(x/(x*(2-x)*(3-x))))**4

epsilon = 1e-8 ### so we don't divide by zero
xmin=0+epsilon
xmax=2-epsilon
N=500000
xlist = []
ylist =[]
ylist_mvm=[]
count = 0
sum = 0
for i in range(N):
  randx = 2*random()
  randy = random()
  val = f(randx)
  sum = sum + val
  if (randy < val):
    count = count + 1
  ylist.append(2*count/(i+1))
  xlist.append(i+1)
  ylist_mvm.append(2*sum/(i+1))


ax = plot.figure(figsize=(10,10))
plot.ylim(0.4,0.42)
plot.scatter(xlist,ylist,color='b',label='MC integration')
plot.scatter(xlist,ylist_mvm,color='r',label='Mean value')
plot.legend(loc='upper right')
plot.ylabel("Estimated value of integral")
plot.xlabel("Trials")
plot.show()
```

Similar performance, though mean value method has slightly better convergence

In many dimensions

One dimension:

$$I = \frac{b-a}{N} \sum_{i=1}^{N} f(x_i)$$

Many dimensions:

$$I = \frac{V}{N} \sum_{i=1}^{N} f(\mathbf{r_i})$$

MC integration works well
in many dimensions,
unlike more traditional
approaches

MC integration to estimate pi

Circle with radius 1 centered at the origin. It has area $\pi r^2 = \pi$

$$x^2 + y^2 = 1$$

Can rewrite as:

$$y = \sqrt{1 - x^2}$$

The area that function sweeps out from x=0 to x=1 is one quadrant of a circle's area. If you do not see this, let's draw on the board

$$\int_0^1 \sqrt{1 - x^2} = \pi/4$$

MC integration to estimate pi


```python
# The "slow" way, maybe easier at first, and still faster than much of the above
import numpy

N = 500000
x_values = numpy.random.uniform(0, 1, N)
results = numpy.zeros(N)

print("done generating random numbers")

for attempt in range(N):
    results[attempt] = math.sqrt(1-x_values[attempt]**2)

# our estimate of the integral (=pi/4) is the total value of results, times the range we're choosing random numbers over
# (1), divided by the number of attempts (= the mean result * 1)
print(f"Our estimate of pi: {results.mean()*4}")
```

```
done generating random numbers
Our estimate of pi: 3.1401902617790585
```

```
[7] # the "fast" way, using numpy tricks
    import numpy
    N = 500000
    x_values = numpy.random.uniform(0, 1, N)
    results = numpy.sqrt(1-x_values**2) # evaluate all together
    print(f"Our estimate of pi: {results.mean()*4}")
```

```
Our estimate of pi: 3.137673840785929
```

Sometimes MC just allows us to be lazy

Problem: What is the probability that 10 dice throws add up exactly to 32?

Exact Way. Calculate this exactly by counting all possible ways of making 32 from 10 dice.

MC method (aka lazy way): Approximate (Lazy) Way. Simulate throwing the dice N times, count the number of times the results add up to 32, and divide this by N

Well, the computer can do the exact way, too

Problem: What is the probability that 10 dice throws add up exactly to 32?

Nicely generates all combinations, with replacement for us!

```python
# Cartestian Product check
from itertools import product

die = [1,2,3,4,5,6]
Ndice = 2

comb = product(die,repeat=Ndice)
for i in list(comb):
  print(i)
```

```
(1, 1)
(1, 2)
(1, 3)
(1, 4)
(1, 5)
(1, 6)
(2, 1)
(2, 2)
(2, 3)
(2, 4)
(2, 5)
(2, 6)
(3, 1)
(3, 2)
(3, 3)
(3, 4)
(3, 5)
(3, 6)
(4, 1)
(4, 2)
(4, 3)
(4, 4)
(4, 5)
(4, 6)
(5, 1)
(5, 2)
(5, 3)
(5, 4)
(5, 5)
(5, 6)
(6, 1)
(6, 2)
(6, 3)
(6, 4)
(6, 5)
(6, 6)
```

Well, the computer can do the exact way, too

Problem: What is the probability that 10 dice throws add up exactly to 32?

Answer: 6.3%

```
# Combinatorics problem
from itertools import product
from math import pow
die = [1,2,3,4,5,6]
Ndice = 10
Ncomb = pow(len(die),Ndice)
toget = 32
nmatch = 0
for i in product(die,repeat=Ndice):
  if (sum(i) == toget): nmatch = nmatch + 1
print("We found ",nmatch," out of ", Ncomb, "for a probability of ",float(nmatch)/Ncomb)
```

```
We found  3801535  out of  60466176.0 for a probability of  0.06287043850763772
```

MC method

Problem: What is the probability that 10 dice throws add up exactly to 32?

MC method Answer: 6.3%

```
# Combinatorics problem
# MC method
import random
Ndice = 10
toget = 32
nmatch = 0
Ntrial = 5000000
for trial in range(Ntrial):
  sum = 0
  for die in range(Ndice):
    roll = random.randint(1,6) ### a single die
    sum = sum+roll
  if (sum == toget): nmatch = nmatch + 1
print("We found ",nmatch," out of ", Ntrial, "for a probability of ",float(nmatch)/Ntrial)
```

```
We found  315030  out of  5000000 for a probability of  0.063006
```

Importance sampling

MC integration will have trouble in regions where the function diverges. And there is anyway no reason why we necessarily need to sample uniformly to perform MC integration - if we can roughly approximate the original function (sample more in regions that "count more") we will have quicker convergence in our results

```
# function to integrate
import matplotlib.pyplot as plt
from numpy import exp,power,frompyfunc,linspace,sqrt,pi

def func(x):
  return (1./sqrt(x))*(1+exp(x))

xmin=1e-5 ##don't start at 0 where function diverges!
xmax=10
N=1000

plt.plot(linspace(xmin,xmax,N), frompyfunc(func, 1, 1)(linspace(xmin,xmax,N)))
plt.show()
```

$$I = \int_0^1 \frac{1}{\sqrt{x}(e^x + 1)} dx$$

Diverges at x = 0!

Importance sampling

$$I = \int_a^b dx f(x) = \int_a^b dx \frac{f(x)}{w(x)} w(x) dx$$

So if we sample not from f but instead from f/w (and multiply by the extra factor of the integral of w), we still have an unbiased estimate

$$I = \frac{1}{N} \sum_{i=1}^{N} \frac{f(x_i)}{w(x_i)} \int_a^b w(x) dx$$

In other words, we can sample points non-uniformly as long as we more heavily weigh points where there are fewer estimates

Importance sampling

$$I = \frac{1}{N} \sum_{i=1}^{N} \frac{f(x_i)}{w(x_i)} \int_a^b w(x) dx$$

If w(x) = 1:

$$I = \frac{1}{N} \sum_{i=1}^{N} \frac{f(x_i)}{1} \int_a^b 1 dx$$

$$I = \frac{1}{N} \sum_{i=1}^{N} f(x_i)(b-a)$$

Get back earlier nominal result from the mean value method!

Importance sampling

```
# Another function for importance sampling
import matplotlib.pyplot as plt
from numpy import exp,power,frompyfunc,linspace,sqrt,pi,random,sin,abs,cos

alpha = 0.01
N=100000
# random integral
def func(x):
  return exp(-x*x*alpha*(abs(sin(x))+abs(cos(x))))

xmin=0
xmax=50
plt.plot(linspace(xmin,xmax,N), frompyfunc(func, 1, 1)(linspace(xmin,xmax,N)))
plt.show()
```

Can also be used with the mean value method to estimate integrals that extend to infinity! Without it we can't do this, since we can't choose random numbers in an infinite uniform range, but we can choose them non-uniformly (ex: Gaussian or exponential)

Importance sampling - nice example from Bryan Webber

$$I = \int_0^1 dx \cos\tfrac{\pi}{2}x = 0.637 \pm 0.308/\sqrt{N}$$

$$I = \int_0^1 dx(1-x^2)\,\frac{\cos\frac{\pi}{2}x}{1-x^2} = \int d\rho\,\frac{\cos\frac{\pi}{2}x}{1-x^2}[x(\rho)] = 0.637 \pm 0.032/\sqrt{N}$$

https://www.hep.phy.cam.ac.uk/theory/webber/MunichPDF/MClecture1.pdf

Stratified Sampling

Divide up integration region piecemeal and optimize to minimize total error.

Can be done automatically (eg VEGAS).

Never as good as Jacobian transformations.

N.B. Puts more points where rapidly varying, not necessarily where larger!

$$I = 0.637 \pm 0.147/\sqrt{N}$$

https://www.hep.phy.cam.ac.uk/theory/webber/MunichPDF/MClecture1.pdf

Multi-dimensional Integration

- Formalism extends trivially to many dimensions
- Particle physics: very many dimensions,
  e.g. phase space = 3 dimensions per particles,
  LHC event ~ 250 hadrons.
- Monte Carlo error remains $\propto 1/\sqrt{N}$
- Trapezium rule $\propto 1/N^{2/d}$
- Simpson's rule $\propto 1/N^{4/d}$

https://www.hep.phy.cam.ac.uk/theory/webber/MunichPDF/MClecture1.pdf

Particle Decays

Simplest example

e.g. top quark decay:

PLUS Integrals over incoming parton choices/densities!

$$|\mathcal{M}|^2 = \frac{1}{2}\left(\frac{8\pi\alpha}{\sin^2\theta_w}\right)^2 \frac{p_t \cdot p_\ell \, p_b \cdot p_\nu}{(m_W^2 - M_W^2)^2 + \Gamma_W^2 M_W^2}$$

$$\Gamma = \frac{1}{2M}\frac{1}{128\pi^3}\int |\mathcal{M}|^2 dm_W^2 \left(1 - \frac{m_W^2}{M^2}\right)\frac{d\Omega}{4\pi}\frac{d\Omega_W}{4\pi}$$

Breit-Wigner peak of W very strong: must be removed by Jacobian factor

https://www.hep.phy.cam.ac.uk/theory/webber/MunichPDF/MClecture1.pdf

Integral without importance sampling


```
# Integral without importance sampling
import matplotlib.pyplot as plt
from numpy import exp,power,frompyfunc,linspace,sqrt,pi,random,sin,abs,cos,array,sum,array

### redo using flat sampling from 0 to 200 (integral rapidly goes to zero)
alpha = 0.01
N=10000000
alpha_inv = 1./alpha

# random integral
def func(x):
  return exp(-x*x*alpha*(abs(sin(x))+abs(cos(x))))

xmin=0
xmax=200
## choose sample alpha as above, but it doesn't have to be!
x_values = random.uniform(0,50,N)
vals = array([func(xi) for xi in x_values])
### now need to multiply by (b-a)
integral = (xmax - xmin)*sum(vals)/N
print("Integral using N = ",N,"points is ",integral)
```

Need a lot of points, and we don't know if we are sampling the tails at all correctly

```
Integral using N =  10000000 points is  31.529579494241684
```

Integral with importance sampling


```
# Integrate with importance sampling
import matplotlib.pyplot as plt
from numpy import exp,power,frompyfunc,linspace,sqrt,pi,random,sin,abs,cos,array,sum,array

alpha = 0.01
N=1000000

# random integral
def func(x):
  return exp(-x*x*alpha*(abs(sin(x))+abs(cos(x))))

## Choose a width of 10 based on above graph, doesn't have to be, of course. We use absolute value since we only care about x>0!
sigma=10
x_values = abs(random.normal(0,sigma,N))

def gauss_func(x):
  return 1./(sqrt(2*pi)*sigma)*exp((-0.5*x*x)/(2*sigma*sigma))


vals = array([func(xi)/gauss_func(xi) for xi in x_values])
## integral was defined conveniently to have unit normalization from 0 to infinity, we used only half but we doubled with absolute value
integral = sum(vals)/N
print("Integral using N = ",N,"points is ",integral)
```

```
Integral using N =  1000000 points is  14.402286823430092
```

Used a Gaussian, but that was just one choice, could have chosen an exponential, or something else that extends to infinity

Statistical mechanics

We need to be clever: even just 500 electrons in naive up/down configurations (only 2 choices!) have 2^{500} possible configurations for the system.

Aurora supercomputer

Intel to Deliver the **Aurora Supercomputer** in 2021

The Argonne National Laboratory **Supercomputer** will accelerate the convergence of high performance computing (HPC) and artificial intelligence with the U.S.' first Exascale system.

www.intel.com › supercomputing › exascale-computing

Aurora Supercomputer - Intel

Exascale = 10^{18} floating point operations / sec

```
>>> 2**500/(10**18)
3.273390607896142e+132
```

```
>>> sec=2**500/(10**18)
>>> min = sec/60.
>>> hour = min/60.
>>> day = hour/24.
>>> year = day/365.
>>> year
1.037985352579953 8e+125
```

We can just sample the system proportionally to the probabilities of the states (based on Boltzmann probabilities, and if we can do that then we can calculate the expectation value for any observable:

$$< X >= \frac{1}{N} \sum_{k=1}^{N} X_k$$

Easy-peasey, right? The problem is that we don't know these probabilities a priori

Markov Chain MC

MCMC: We start the system out in some set of states. And then we generate a new set of states from a fixed set of probabilistic rules, depending only on the previous state. And another one from the same set of rules, again depending only on the most recent, previous state. This fully collection or chain of states is the Markov Chain.

The probabilities rules are the transition probabilities T_{ij}, which define the probability of changing from state i to state j. We will define them so that the probability of being at any one state on the chain is the Boltzmann probability

$$\sum_j T_{ij} = 1$$

Sum over all possible transition probabilities (which includes probabilities for no transmission!) must be 1

How to choose transition probabilities?

$$\frac{T_{ij}}{T_{ji}} = \frac{P(E_j)}{P(E_i)} = \frac{e^{-\beta E_j}/Z}{e^{-\beta E_i}/Z} = e^{-\beta(E_j - E_i)}$$

Look at the ratio of probability to go from state i to state j to the probability to go backwards, from state j to state i. Set this to the ratio of probabilities of being in those states from Boltzmann. Note that the Z above is the tricky thing that we don't typically know how to calculate, but it cancels out in the ratio

We are not going to start initially in a correct state for the system according to Bolzmann (we don't know it in advance), but what happens if we do reach that state in the ith step on the chain? Where do we get to in the jth step?

How to choose transition probabilities?

$$\frac{T_{ij}}{T_{ji}} = \frac{P(E_j)}{P(E_i)} = \frac{e^{-\beta E_j}/Z}{e^{-\beta E_i}/Z} = e^{-\beta(E_j - E_i)}$$

We are not going to start initially in a correct state for the system according to Bolzmann (we don't know it in advance), but what happens if we do reach that state in the ith step on the chain? Where do we get to in the jth step?

$$\sum_i T_{ij} P(E_i) = \sum_i T_{ji} P(E_j) = P(E_j) \sum_i T_{ji} = P(E_j)$$

So in the jth step, the probability of being in any state is also given correctly. And then in the kth step. So once we get to a Boltzmann distribution we stay there! (Need to run the system long enough to converge to it)

Metropolis Algorithm

If we are in ith state and want to know whether to move to a jth state, we do so with some probability. If the move decreases the energy of the system, we make that choice. If it doesn't, we still may allow it with acceptance probability Pa

$$P_a = 1, \; E_j \leq E_i$$

$$P_a = e^{-\beta(E_j - E_i)}, \; E_j > E_i$$

Of course, that is the probability of making a specific, proposed move give, we need to know the probability of making that move vs all other moves. Following the book, consider the case that there are M possible moves

Metropolis Algorithm

$$P_a = 1, \; E_j \leq E_i$$

$$P_a = e^{-\beta(E_j - E_i)}, \; E_j > E_i$$

Of course, that is the probability of making a specific, proposed move give, we need to know the probability of making that move vs all other moves. Following the book, consider the case that there are M possible moves:

$$T_{ij} = \frac{1}{M} e^{-\beta(E_j - E_i)}, \; E_j > E_i$$

$$T_{ji} = \frac{1}{M}, \; E_j > E_i$$

We always accept move back since it lowers energy here

Metropolis Algorithm

$$T_{ij} = \frac{1}{M} e^{-\beta(E_j - E_i)}, \; E_j > E_i$$

$$T_{ji} = \frac{1}{M}, \; E_j > E_i$$

$$\frac{T_{ij}}{T_{ji}} = \frac{\frac{1}{M} e^{-\beta(E_j - E_i)}}{1/M}, \; E_j > E_i$$

$$\frac{T_{ij}}{T_{ji}} = e^{-\beta(E_j - E_i)}, \; E_j > E_i$$

Just what we want!!!

If $E_i < E_j$ the same holds true! (Will leave it to you to double-check)

Metropolis Algorithm

So how do we use this?

1) We start at some random state

2) Consider the move to one new state from the allowed set of moves

3) We calculate the acceptance probability and use that to decide whether to make the move or not. If we accept the move, go ahead and change the state

4) If we reject the move, stay in the current state (but count the step!)

5) Continue to monitor the quantity of interest and see if it stabilizes/equilibrates

Assumes the system is ergodic, meaning that you can get to every possible state!

Ising model

Set of N spins on a lattice, assume they can have only two values: up or down

Force between magnets falls rapidly with distance, so for any given magnet consider only its "nearest neighbors"

$s_i = +1$ (spin up) or
-1 (spin down)

Ising model

s_i = +1 (spin up) or
-1 (spin down)

$$M = \sum_i s_i$$

$$E = -J \sum_{i,j} s_i s_j - H \sum_i s_i$$

Spins will align to H

Sum ONLY over
"nearest neighbors"

H is an external field
that couples to the
total magnetization M

J>0: ferromagnetic material, energy minimized if spins aligned
J<0: anti-ferromagnetic material, energy minimized if spins
locally point in opposite direction

Ising model

↓	↓	↑	↑	↑	↓	↑	↓	↑	↑
↓	↓	↑	↓	↑	↑	↓	↑	↓	↓
↓	↑	↓	↑	↓	↑	↓	↓	↑	↑
↓	↑	↓	↓	↑	↓	↓	↑	↑	↓
↑	↓	↑	↓	↓	↓	↑	↑	↓	↑
↑	↑	↓	↓	↓	↓	↑	↑	↑	↓
↓	↓	↑	↓	↑	↓	↑	↓	↑	↓
↑	↓	↓	↑	↑	↓	↓	↑	↑	↓
↓	↓	↑	↑	↑	↓	↑	↓	↑	↓
↓	↓	↑	↓	↓	↑	↓	↑	↑	↑

For H=0, the system will be in one of two states
low temperature (below Curie temperature): magnetized
high temperature (above Curie temp): sum of magnetization zero

Ising model tricks

↓	↓	↑	↑	↑	↓	↑	↓	↑	↑
↓	↓	↑	↓	↑	↑	↓	↑	↓	↓
↓	↑	↓	↑	↓	↑	↓	↓	↑	↑
↓	↑	↓	↓	↑	↓	↓	↑	↑	↓
↑	↓	↑	↓	↓	↓	↑	↑	↓	↑
↑	↑	↓	↓	↓	↓	↑	↑	↑	↓
↓	↓	↑	↓	↑	↓	↑	↓	↑	↓
↑	↓	↓	↑	↑	↓	↓	↑	↑	↓
↓	↓	↑	↑	↑	↓	↑	↓	↑	↓
↓	↓	↑	↓	↓	↑	↓	↑	↑	↑

For fixed T and H==0, only have a finite number (five!) of possible values for the sum of neighboring spins!
And for H !=0, only have 10 total possible values

Ising model tricks

↓	↓	↑	↑	↑	↓	↑	↓	↑	↑
↓	↓	↑	↓	↑	↑	↓	↑	↓	↓
↓	↑	↓	↑	↓	↑	↓	↓	↑	↑
↓	↑	↓	↓	↑	↓	↓	↑	↑	↓
↑	↓	↑	↓	↓	↓	↑	↑	↓	↑
↑	↑	↓	↓	↓	↓	↑	↑	↑	↓
↓	↓	↑	↓	↑	↓	↑	↓	↑	↓
↑	↓	↓	↑	↑	↓	↓	↑	↑	↓
↓	↓	↑	↑	↑	↓	↑	↓	↑	↓
↓	↓	↑	↓	↓	↑	↓	↑	↑	↑

Important to take advantage of numpy and parallelization.
Assume periodic boundary conditions at the “edges” of the lattice

Statistical Physics and definitions

The Boltzmann distribution is given by: $$p(\mathscr{C}) = \frac{1}{Z_N} \exp\left[-\frac{E(\mathscr{C})}{k_B T}\right]$$

where E($\mathcal{C}$) is the energy of a spin configuration $\mathcal{C}$

The partition function is defined as

$$Z_N = \sum_{\mathscr{C}} \exp\left[-\frac{E(\mathscr{C})}{k_B T}\right]$$

from which we can derive the average energy

$$\langle E \rangle = \sum_{\mathscr{C}} p(\mathscr{C}) E(\mathscr{C}) = k_B T^2 \frac{\partial}{\partial T} \ln Z_N$$

and magnetization

$$\langle M \rangle = \sum_{\mathscr{C}} p(\mathscr{C}) \mathscr{M}(\mathscr{C}) = k_B T \frac{\partial}{\partial h} \ln Z_N$$

using these we can define

$$\chi = \frac{\partial}{\partial h} \langle M \rangle$$ **susceptibility**

$$c_h = \frac{\partial}{\partial T} \langle E \rangle$$ **specific heat**

From Andreas, a nice summary

Final expressions are:

$$c_h = \frac{1}{k_B} T^2 \sum_{\mathscr{C}} p(\mathscr{C}) \left[E^2(\mathscr{C}) - E(\mathscr{C}) \langle E \rangle \right]$$

$$= \frac{1}{k_B T^2} \left(\langle E^2 \rangle - \langle E \rangle^2 \right)$$

$$= \frac{1}{k_B T^2} \mathrm{var}\,(E) \ .$$

$$\chi = \frac{1}{k_B T} \sum_{\mathscr{C}} p(\mathscr{C}) \left[\mathscr{M}^2(\mathscr{C}) - \mathscr{M}(\mathscr{C}) \langle M \rangle \right]$$

$$= \frac{1}{k_B T} \left(\langle M^2 \rangle - \langle M \rangle^2 \right)$$

$$= \frac{1}{k_B T} \mathrm{var}\,(M) \ .$$

1D solution

(see book for details)

The partition function in 1D for N spins can be calculated as

$$Z_N = \lambda_1^N + \lambda_2^N$$

with

$$\lambda_{1,2} = \exp\left(\frac{J}{k_B T}\right)\cosh\left(\frac{h}{k_B T}\right) \pm \sqrt{\exp\left(\frac{2J}{k_B T}\right)\sinh^2\left(\frac{h}{k_B T}\right) + \exp\left(-\frac{2J}{k_B T}\right)}$$

The expectation value for the energy per particle is given by

$$\langle \varepsilon \rangle = \frac{k_B T^2}{N}\frac{\partial}{\partial T}\ln Z_N$$

In the thermodynamic limit N→∞:

$$\lim_{N\to\infty}\frac{1}{N}\cancel{Z_N} = \lim_{N\to\infty}\frac{1}{N}\ln\left(\lambda_1^N + \lambda_2^N\right) = \ln\lambda_1$$

$\ln(Z_N)$

and for h=0:

$$\boxed{\lim_{N\to\infty}\frac{1}{N}\cancel{Z_N} = \ln\left[2\cosh\left(\frac{J}{k_B T}\right)\right]}$$

smooth function of T for T > 0
→ no phase transition in the one dimensional Ising model

2D Onsager solution

For h=0 one observes a second order phase transition with transition temperature defined by:

c_h and χ diverge at T_c

$$2\tanh^2\left(\frac{2J}{k_B T_C}\right) = 1 \qquad k_B T_c = \frac{2J}{\log(1+\sqrt{2})} \approx 2.269J$$

and for the energy per particle

$$\langle\varepsilon\rangle = -J\coth\left(\frac{2J}{k_B T}\right)\left\{1 + \frac{2}{\pi}K_1(\xi)\left[2\tanh^2\left(\frac{2J}{k_B T}\right) - 1\right]\right\}$$

where $K_1(\xi)$ is the complete elliptic integral with

$$\xi = \frac{2\sinh\left(\frac{2J}{k_B T}\right)}{\cosh^2\left(\frac{2J}{k_B T}\right)}$$

and magnetization per particle

$$\langle m\rangle = \begin{cases} \dfrac{(1+z^2)^{\frac{1}{4}}(1-6z^2+z^4)^{\frac{1}{8}}}{\sqrt{1-z^2}} & \text{for } T < T_C \\ 0 & \text{for } T > T_C \end{cases}$$

with $z = \exp\left(-\frac{2J}{k_B T}\right)$

$$= \left(1 - \left[\sinh\left(\log(1+\sqrt{2})\frac{T_c}{T}\right)\right]^{-4}\right)^{\frac{1}{8}} \qquad T<T_c$$

Phase transitions

Ehrenfest classification of Phase Transition:

- **First-order phase transitions** exhibit a discontinuity in the first derivative of the chemical potential with a thermodynamic variable. Such as solid/liquid/gas transitions.
- **Second-order phase transitions** (also called continuous phase transition) have a discontinuity or divergence in a second derivative of the chemical potential with thermodynamic variables.

c_h and χ are second derivatives

Critical exponents

Reduced temperature: $\tau \equiv \frac{T - T_C}{T_C}$

Critical exponent: $k \stackrel{\text{def}}{=} \lim_{\tau \to 0} \frac{\log |f(\tau)|}{\log |\tau|}$

Specific heat $C \propto |\tau|^{-\alpha}$ ($T<T_c$)

Magnetization $M \propto |\tau|^{\beta}$

Magnetic susceptibility $\chi \propto |\tau|^{-\gamma}$

Correlation length $\xi \propto |\tau|^{-\nu}$

Critical behavior of the order parameter an the correlation length. The order parameter vanishes with the power β of the reduced temperature t as the critical point is approached along the line of phase coexistance. The correlation length diverges with the power ν of the reduced temperature.

The exponents display critical point universality (don't depend on details of the model). This explains the success of the Ising model in providing a quantitative description of real magnets.

Ising values

d	2	3	4
α	0 (log div)	0.110(1)	0
β	1/8	0.3265(3)	1/2
γ	7/4	1.2372(5)	1
δ	15	4.789(2)	3
η	1/4	0.0364(5)	
ν	1	0.6301(4)	1/2
ω	2	0.84(4)	

Ising model


```
# Ising model 10.9
# using numpy array manipulation. Note that we calculate the energy sum over and over,
# so even though numpy is efficient, the overall calculation is not
from math import exp
from numpy import empty,sum,arange
from random import random, randrange
from pylab import plot,show,xlabel,ylabel,subplots

L = 50
N = 10000000
J = 1.6
T = 2.3

# Function to calculate the energy
def energy(s):
  return -J*(sum(s[0:L-1,:]*s[1:L,:]) + sum(s[:,0:L-1]*s[:,1:L]))

# Initial state
s = empty([L,L],int)
for i in range(L):
  for j in range(L):
    if random() < 0.5:
      s[i,j] = +1
    else:
      s[i,j] = -1
E = energy(s)
M = sum(s)
```

Ising model

```
# Main loop
mpoints = []
epoints = []
xpoints = arange(N)
for k in range(N):
  mpoints.append(M)
  epoints.append(E)

  # Save current energy
  oldE = E

  # Choose a random spin, flip it, and calculate dE.
  i = randrange(L)
  j = randrange(L)
  s[i,j] = -s[i,j]
  # We could probably be move clever since "most" of the calculation doesn't change, only nearest neighbors
  E = energy(s)
  deltaE = E - oldE

  # Decide whether to accept the move or not
  if deltaE > 0.0: ### If dE < 0 we always keep things
    if random() > exp(-deltaE/T):
      # Move rejected, revert to old state, don't need to recalculate M, we haven't changed it
      s[i,j] = -s[i,j]
      E = oldE
      continue
  # Accepted! Calculate new values
  M = sum(s)
```

Let’s discuss the results!

Ising model tricks

Can see the magnetization quickly approaching the maximum value it can have (in absolute value). In another iteration that could be positive, not negative. But now let's try and be smarter about how we calculate things (also being more careful about edge effects), though we still use numpy initially and also to calculate the magnetization

```
from math import exp
from numpy import empty,sum,arange
from random import random, randrange
from pylab import plot,show,xlabel,ylabel,subplots

L = 50
N = 10000000
J = 1.6
T = 2.3

# Function to calculate the energy
def energy(s):
  return -J*(sum(s[0:L-1,:]*s[1:L,:]) + sum(s[:,0:L-1]*s[:,1:L]))

# Initial state
s = empty([L,L],int)
for i in range(L):
  for j in range(L):
    if random() < 0.5:
      s[i,j] = +1
    else:
      s[i,j] = -1
E = energy(s)
M = sum(s)
```

Ising model tricks


```
# Main loop
mpoints = []
epoints = []
xpoints = arange(N)
for k in range(N):
  mpoints.append(M)
  epoints.append(E)

  # Save current energy
  oldE = E

  # Choose a random spin, flip it, and calculate dE
  i = randrange(L)
  j = randrange(L)
  # We could probably be move clever since "most" of the calculation doesn't change, only nearest neighbors
  # So let's try that here. We will flip the spin AFTER after calculating dE, too
  #E = energy(s)
  #deltaE = E - oldE
  iup = i+1
  idown = i-1
  jup = j+1
  jdown = j-1
  if (iup == L): iup = 0
  if (idown == 0): idown = L-1
  if (jup == L): jup = 0
  if (jdown == 0): jdown = L-1
  ### Factor of two comes from (1 - (-1) = 2)
  deltaE = 2*J*s[i,j]*(s[iup,j]+s[idown,j]+s[i,jup]+s[i,jdown])
  E = deltaE + oldE
  ### Flip!
  s[i,j] = -s[i,j]
```

Ising model tricks

```
# Make the plots


fig, ax1 = subplots()
color = 'tab:red'
ax1.set_xlabel('N')
ax1.set_ylabel('Magnetization', color=color)
ax1.plot(xpoints, mpoints, color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()  # instantiate a second axes that shares the same x-axis

color = 'tab:blue'
ax2.set_ylabel('E', color=color)  # we already handled the x-label with ax1
ax2.plot(xpoints, epoints, color=color)
ax2.tick_params(axis='y', labelcolor=color)

fig.tight_layout()  # otherwise the right y-label is slightly clipped
```

Ising model, bigger

That was faster! Let's try something bigger, so we use more N, too

```
[ ] from math import exp
    from numpy import empty,sum,arange
    from random import random, randrange
    from pylab import plot,show,xlabel,ylabel,subplots

    L = 80
    N = 40000000
    J = 1.6
    T = 2.0
```

10.7

Also, consider a 50x50 Ising model with J = 1.0. Make pretty plots of the final set of spin configuration for T = k*Tc where k is a range of values between 0.5 and 2.0 (at least 5 such values). For each configuration, find the energy and magnetization and then plot the magnetization and energy vs T_c. Then do the same with an external magnetic field H = 1.0.

Discuss your results

Final assignment

5.20, 6.16

Also (for everyone!), consider the double pendulum with $L = m_1 = m_2 = 1$ and $\Theta_1(t=0) = 2.0$, $\Theta_2(t=0) = 2.0$, $\omega_1(t=0) = 0.3$. Plot the final Θ_1 and Θ_2 at t = 50s for 100 steps of $\omega_2(t=0)$ from -3 to +3. Be careful about plotting angles (ie modulus of 2π), and also be careful about how you obtain these results (ie number of steps, method used, etc).

Compare this to the single non-linear pendulum with $\Theta(t=0) = 2.0$ and $\omega(t=0)$ from -3 to +3, and discuss your results

Finally, 8.8 (PHYS510 only)